

Deep Learning

Indrajith Wasala Mudiyanse

Section 01

Question 01

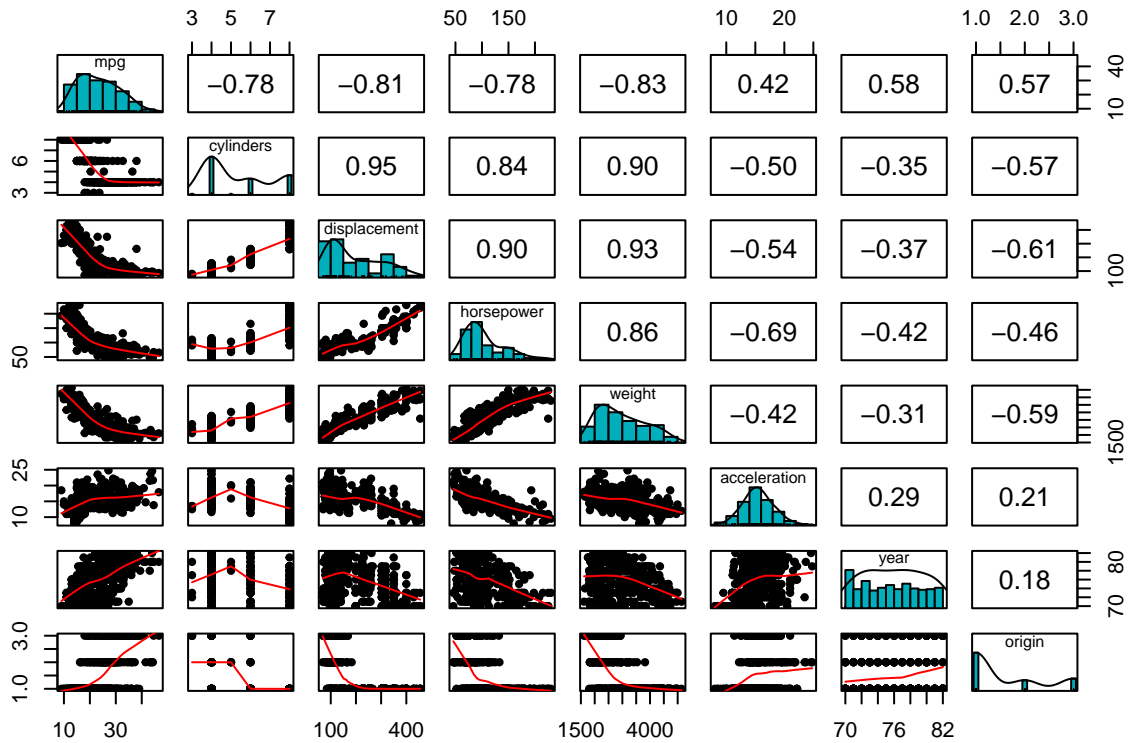


Figure 1: Pairwise scatterplots

- a) The pairwise scatter plot visually reveals that the suggested response variable (mpg) has a strong negative linear relationship with the variables cylinders, displacement, horsepower and weight and moderately positive linear relationship between acceleration, year and origin. Also, note that predictor variables cylinders, displacement, horsepower and weight have a strong positive linear relationship between each other while most of the other predictors have weak negative linear relationships. The correlation matrix confirms the above information. The Histograms visually suggest that the response variable mpg have a right skewed distribution. Also visually, all the predictors does not have symmetric distributions except acceleration.

- b) MSE via LOOCV = 11.37113. Regression coefficients,

##	Estimate	Std. Error	t value	Pr(> t)
## (Intercept)	23.4459184	0.1680733	139.498139	0.000000e+00
## cylinders	-0.8415931	0.5514496	-1.526147	1.277965e-01
## displacement	2.0819598	0.7864080	2.647430	8.444649e-03
## horsepower	-0.6524692	0.5306734	-1.229512	2.196328e-01
## weight	-5.4990690	0.5538510	-9.928787	7.874953e-21
## acceleration	0.2222978	0.2726998	0.815174	4.154780e-01
## year	2.7656487	0.1877716	14.728795	3.055983e-39
## origin	1.1487821	0.2240437	5.127492	4.665681e-07

- c) MSE via LOOCV = 11.37113.

- d) Shallow learning model MSE via LOOCV 488.5207.

	Coefficients
(Intercept)	0.208653
cylinders	-0.49084339
displacement	0.39686212
horsepower	0.15018430
weight	-1.02059686
acceleration	0.60515058
year	-0.05495063
origin	-0.22021827

Table 1: Regression coefficients by shallow learning model

- e) The estimates reported in above parts are not identical. If we run more epochs we may get more closer estimates (The neural network needs more training).

Question 02

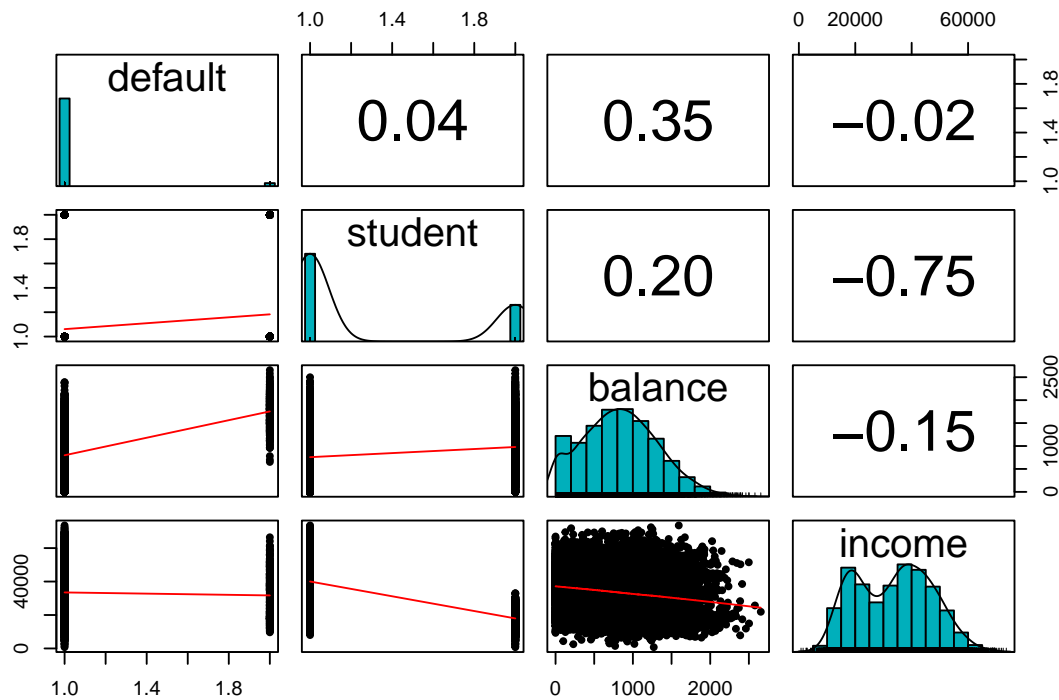


Figure 2: Pairwise scatterplots

- a) The pairwise scatter plot visually reveals that the suggested response variable (default) has a weak positive linear relationship with the variables student and balance while income has weak negative linear relationship. The correlation matrix confirms the above information. The Histograms visually suggest that none of the variables have symmetric distribution.
- b) 5-fold CV estimate of test error is 0.02680055. Logistic regression coefficients (Note that this is without standardizing training data),

##	Estimate	Std. Error	z value	Pr(> z)
## (Intercept)	-1.086905e+01	4.922555e-01	-22.080088	4.911280e-108
## studentYes	-6.467758e-01	2.362525e-01	-2.737646	6.188063e-03
## balance	5.736505e-03	2.318945e-04	24.737563	4.219578e-135
## income	3.033450e-06	8.202615e-06	0.369815	7.115203e-01

- c) 5-fold CV estimate of test error is 0.0268.
- d) 5-fold CV test error of the shallow learning model is 0.03545891.

	Coefficients
(Intercept)	-2.537609
studentYes	-5.6365714073
balance	0.0167175867
income	-0.0008408058

Table 2: Regression coefficients by shallow learning model

- e) The estimates reported in above parts are not identical. If we run more epochs we may get more closer estimates (The shallow learning model needs more training).

Question 03

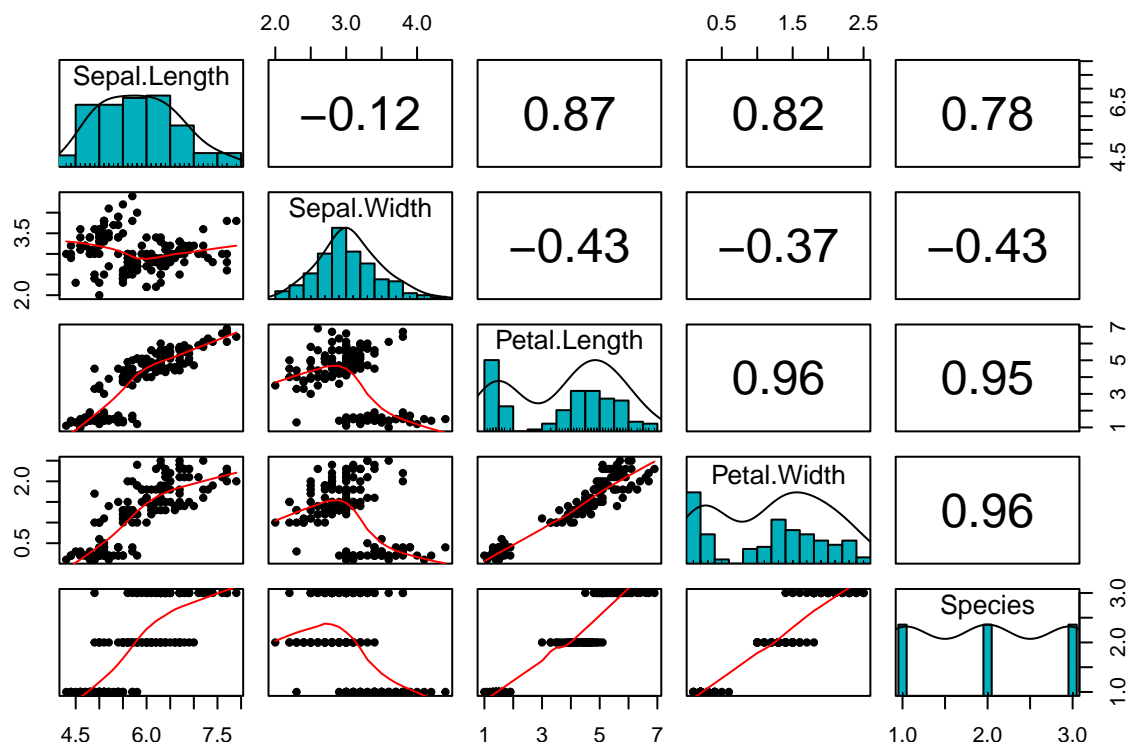


Figure 3: Pairwise scatterplots

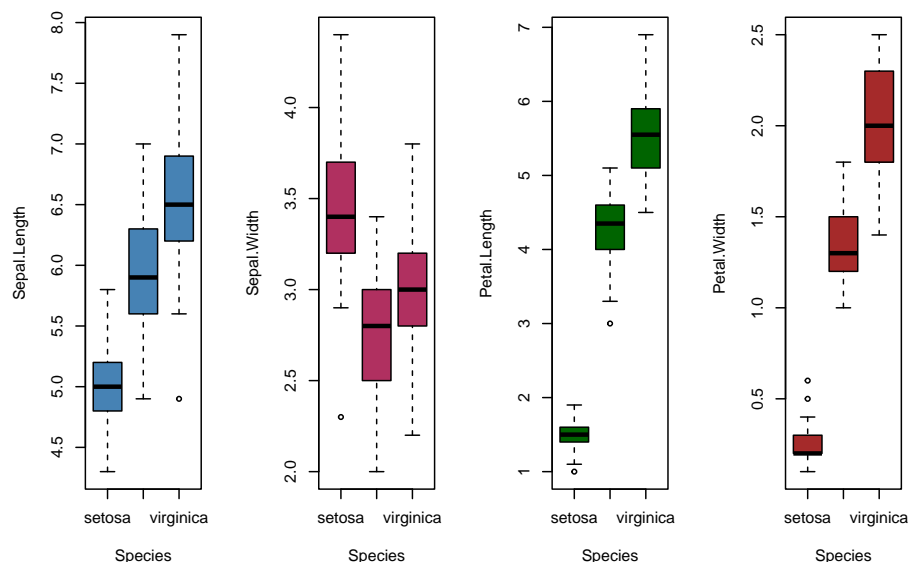


Figure 4: Boxplots of response variable with predictor variables

- a) The pairwise scatter plot visually reveals that the suggested response variable (Species) has a strong positive linear relationship with all the variables except Sepal.Width. The correlation matrix confirms the above information. The Histograms visually suggest that non of the variables have symmetric distribution except Sepal.Width. Moreover the side by side boxplots reveal the distribution of the predictors according to the levels of the response variable.

- b) 5-fold CV estimate of test error is 0.0266667. The regression coefficients,

	(Intercept)	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
## versicolor	9.023056	-3.345346	-3.039437	13.69517	8.090794
## virginica	-11.021851	-5.380110	-5.945487	30.29748	21.980253

- c) 5-fold CV test error of the shallow learning model is $1 - 0.6933334 = 0.3066666$.

		Coefficients	
(Intercept)	-0.267862678	0.199472770	-0.007312237
Sepal.Length	-0.6619897	0.3302517	0.6933144
Sepal.Width	-0.5235697	0.1453871	-0.1324961
Petal.Length	-0.7744051	1.1369798	1.2270440
Petal.Width	-0.2853383	-0.6473362	-0.2898701

Table 3: Regression coefficients by shallow learning model

- d) The estimates reported in above parts are not identical. If we run more epochs we may get more closer estimates (The shallow learning model needs more training).

Section 2 (R codes)

```
knitr::opts_chunk$set(echo = TRUE}  
  
## ----echo=FALSE,warning=FALSE,fig.align="center",fig.cap="Pairwise scatterplots",out.width = "100%"----  
  
### Question 01  
  
## a)  
  
library(ISLR)  
library(psych)  
pairs.panels(Auto[,-9],  
             method = "pearson", # correlation method  
             hist.col = "#00AFBB",  
             density = TRUE, # show density plots  
             ellipses = FALSE # show correlation ellipses  
             )  
  
## ----echo=FALSE-----  
  
## b)  
  
mean <- apply(Auto[,-c(1,9)], 2, mean)  
std <- apply(Auto[,-c(1,9)], 2, sd)  
  
# Standardize the training and test features  
Auto_data <- as.data.frame(scale(Auto[,-c(1,9)], center = mean, scale = std))  
Auto_data1<-cbind(Auto[,1],Auto_data)  
colnames(Auto_data1)[1]<-c("mpg")  
reg<-glm(mpg~., data=Auto_data1) # Fitting regression  
summary(reg)$coefficient  
  
## ----include=FALSE-----  
library(boot)  
cv.est <- cv.glm(Auto_data1, reg)$delta[1]  
  
## ----include=FALSE-----  
  
## c)  
  
LOOCV<-function(dataset){  
  n<-length(dataset[,1])  
  RSS_i<-c()  
  for (i in 1:n) {  
    newdata<-dataset[-i,] # Data without ith observation (Training data)  
    testdata<-dataset[i,] # ith observation  
    fit <- glm(mpg ~ ., data = newdata) # Fit linear regression for Training data  
    lr.fit <- predict(fit, testdata, type = "response")  
    RSS_i[i] <- (testdata$mpg-lr.fit)^2 # Get RSS for test  
  }  
  MSE<- mean(RSS_i)  
  return(list(LOOCV_MSE=MSE))  
}  
LOOCV(dataset=Auto_data1)  
  
## ---- include=FALSE-----  
  
## d)
```

```

mean <- apply(Auto[, -c(1,9)], 2, mean)
std <- apply(Auto[, -c(1,9)], 2, sd)
# Standardize the training and test features
Auto_data <- as.data.frame(scale(Auto[, -c(1,9)], center = mean, scale = std))
Auto_data1 <- cbind(Auto[, 1], Auto_data)
colnames(Auto_data1)[1] <- c("mpg")

train_data <- as.matrix(Auto_data1[, -1])
train_targets <- as.matrix(Auto_data1[, 1], ncol=1)

# Specify and Pre compile the model
build_model <- function(){
  # specify the model
  model <- keras_model_sequential() %>%
  layer_dense(units = 1, input_shape = dim(train_data)[2], activation = 'linear')
  # compile the model
  model %>% compile(
    optimizer = "rmsprop",
    loss = "mse")
}

MSE <- c()

for (i in 1:nrow(train_data)) {
  val_data <- matrix(train_data[i,], nrow=1)
  val_targets <- train_targets[i]
  partial_train_data <- train_data[-i,]
  partial_train_targets <- train_targets[-i]

  model <- build_model() # use precompiled model function

  model %>% fit(partial_train_data, partial_train_targets,
    epochs = 50, batch_size = 128, verbose = 0) # trains the model in silent mode (verbose=0)
  pred_targets <- predict(model, val_data)
  MSE[i] <- (val_targets - pred_targets)^2
  # weights <- model %>% get_layer(index = 1) %>% get_weights()
}

weights <- model %>% get_layer(index = 1) %>% get_weights()

#####
## ----echo=FALSE, warning=FALSE, fig.align="center", fig.cap="Pairwise scatterplots", out.width = "80%"----

### Question 02

## a)

library(ISLR)
library(psych)
pairs.panels(Default,
  method = "pearson", # correlation method
  hist.col = "#00AFBB",
  density = TRUE, # show density plots
  ellipses = FALSE # show correlation ellipses
)

## ----echo=FALSE-----

## b)

```

```

logi_reg <- glm(default ~ ., data = Default, family = "binomial")
summary(logi_reg)$coefficient

## ----include=FALSE-----
# Using caret package

library(caret)

glm.fit2 <- train(
  form = default ~ .,
  data = Default,
  trControl = trainControl(method = "cv", number = 5,
    seeds = set.seed(1)),
  method = "glm",
  family = "binomial")

Miss_clas<-1-glm.fit2$results[[2]] # Misclassification rate = 1 - Accuracy,

## ----include=FALSE-----

## c)

MY_CV<-function(dataset,fold=5){

  k<-fold
  set.seed(1)
  indices <- sample(1:nrow(dataset),replace = FALSE)
  folds <- cut(indices, breaks = k, labels = FALSE)
  pred<-c()
  lr.pred.fit<-c()

  for (i in 1:k) {
    val_indices <- which(folds == i, arr.ind = TRUE) # prepares the validation data: data from partition #k
    val_data <- Default[val_indices,]
    partial_train_data <- Default[-val_indices,]
    fit <- glm(default ~ ., family = binomial, data = partial_train_data)
    lr.prob.fit <- predict(fit, val_data, type = "response")
    pred <- ifelse(lr.prob.fit >= 0.5, "Yes", "No")
    lr.pred.fit[i]<-mean(pred==val_data[, "default"])
  }

  MSE<- 1 - mean(lr.pred.fit)

  return(list(MSE=MSE))
}
Error<-MY_CV(dataset=Default)

## ---- include=FALSE-----

## d)

train_data<-Default[,-1] # Without default
train_data[,1]<-as.numeric(unclass(train_data[,1])-1) # Convert student (No=0,Yes=1)
train_data<-as.matrix(train_data)
train_targets<-as.matrix((unclass(Default[,1])-1),ncol=1) # Converted default (No=0,Yes=1)

build_model <- function(){
  # specify the model
  model <- keras_model_sequential() %>%

```

```

layer_dense(units = 1, input_shape = dim(train_data)[2], activation = 'sigmoid')
# compile the model
model %>% compile(
  optimizer = "rmsprop",
  loss = "binary_crossentropy",
  metrics = c("mae") # mean absolute error
)
}

# K-fold CV
k <- 5
indices <- sample(1:nrow(train_data))
folds <- cut(indices, breaks = k, labels = FALSE)
num_epochs <- 50 # number of epochs
all_scores <- c()
for (i in 1:k){
  cat("Processing fold #", i, "\n")

  val_indices <- which(folds == i, arr.ind = TRUE) # prepares the validation data: data from partition #k
  val_data <- train_data[val_indices,]
  val_targets <- train_targets[val_indices]

  partial_train_data <- train_data[-val_indices,] # prepares the training data: data from all other partitions
  partial_train_targets <- train_targets[-val_indices]

  model <- build_model() # use precompiled model function

  model %>% fit(partial_train_data, partial_train_targets,
    epochs = num_epochs, batch_size = 16,
    verbose = 0) # trains the model in silent mode (verbose = 0)

  results <- model %>% evaluate(val_data, val_targets, verbose = 0) # evaluate model on the validation data
  all_scores <- c(all_scores, results["mae"])
}
all_scores
mean(all_scores)

## ---- include=FALSE-----
weights2 <- model %>% get_layer(index = 1) %>% get_weights()

#####

## ----echo=FALSE,warning=FALSE,fig.align="center",fig.cap="Pairwise scatterplots",out.width = "80%"----

### Question 03

## a)

pairs.panels(iris,
  method = "pearson", # correlation method
  hist.col = "#00AFBB",
  density = TRUE, # show density plots
  ellipses = FALSE # show correlation ellipses
)

## ---- echo=FALSE,warning=FALSE,fig.align="center",fig.cap="Boxplots of response variable with predictor variables"
par(mfrow = c(1, 4))
boxplot(iris$Sepal.Length ~ iris$Species,
  col='steelblue',
  xlab='Species',

```



```

        ylab='Sepal.Length')
boxplot(iris$Sepal.Width ~ iris$Species,
        col='maroon',
        xlab='Species',
        ylab='Sepal.Width')
boxplot(iris$Petal.Length ~ iris$Species,
        col='darkgreen',
        xlab='Species',
        ylab='Petal.Length')
boxplot(iris$Petal.Width ~ iris$Species,
        col='brown',
        xlab='Species',
        ylab='Petal.Width')

## ----include=FALSE-----

## b)

# Standardize the training and test features
mean <- apply(iris[,-c(5)], 2, mean)
std <- apply(iris[,-c(5)], 2, sd)

iris_data <- as.data.frame(scale(iris[,-c(5)], center = mean, scale = std))
iris_data1<-cbind(iris_data,iris[,5])
colnames(iris_data1)[5]<-c("Species")

## ----include=FALSE-----
library(nnet)
mul_reg <- multinom(Species ~ ., data = iris_data1)

## ----echo=FALSE-----
summary(mul_reg)$coefficien

## ----include=FALSE-----
multi.fit <- train(
  form = Species ~ .,
  data = iris_data1,
  trControl = trainControl(method = "cv", number = 5,
    seeds = set.seed(1)),
  method = "multinom")

Miss_clas<-1-multi.fit$results[[2]] # Misclassification rate = 1 - Accuracy,
Miss_clas

## ---- include=FALSE-----

## c)

train_data<-iris_data1[,-5]
train_data<-as.matrix(train_data)
train_targets<-as.matrix((unclass(iris[,5])-1),ncol=1)

train_labels <- to_categorical(train_targets)

build_model <- function(){
  # specify the model
  model <- keras_model_sequential() %>%
    layer_dense(units = 3, input_shape = dim(train_data)[2],activation = 'softmax')

```

```

# compile the model
model %>% compile(
  optimizer = "rmsprop",
  loss = "categorical_crossentropy",
  metrics = c("accuracy") # monitor classification accuracy
)
}

# K-fold CV
k <- 5
indices <- sample(1:nrow(train_data))
folds <- cut(indices, breaks = k, labels = FALSE)
num_epochs <- 100 # number of epochs
all_scores <- c()
for (i in 1:k){
  cat("Processing fold #", i, "\n")

  val_indices <- which(folds == i, arr.ind = TRUE) # prepares the validation data: data from partition #k
  val_data <- train_data[val_indices,]
  val_targets <- train_labels[val_indices,]

  partial_train_data <- train_data[-val_indices,] # prepares the training data: data from all other partitions
  partial_train_targets <- train_labels[-val_indices,]

  model <- build_model() # use precompiled model function

  model %>% fit(partial_train_data, partial_train_targets,
    epochs = num_epochs, batch_size = 32,
    verbose = 0) # trains the model in silent mode (verbose = 0)

  results <- model %>% evaluate(val_data, val_targets, verbose = 0) # evaluate model on the validation data
  all_scores <- c(all_scores, results[2])
}
all_scores
mean(all_scores)

## ---- include=FALSE-----
weights3 <- model %>% get_layer(index = 1) %>% get_weights()

```