

Deep Learning

Indrajith Wasala Mudiyanse

Section 01

Question 01

- a) Please find the code in Section 02. When comparing models with dictionary sizes 1000 and 5000, there is a small difference in accuracy (both are quite similar to each other).

Dictionary size	RNN model accuracy
1000	0.50408
5000	0.50440

Table 1: Test accuracies for RNN model described in Section 10.9.6

- b) As an improvement I have added two extra layers to the above model with ReLu activation (Please check the code in Section 02). After this modification there is an improvement in terms of accuracy.

Dictionary size	RNN model accuracy
1000	0.61397
5000	0.61404
10000	0.67408

Table 2: Test accuracies for improved RNN model described in Section 10.9.6

Question 02

- a) Please check the code in Section 02. After the inclusion of variable “day_of_week” there is a quite large improvement in the model R^2 value.

Model	R^2 value
Original model	0.408254
After adding "day of week"	0.980364

Table 3: Test R^2 value for fitted RNN model

- b) I added a 32-unit dense layer with ReLu activation. The model seems to be slightly improved (Very small improvement. Check the code in Section 02).

Model	R^2 value
Original model	0.408254
Improved model	0.408319

Table 4: Test R^2 value for fitted RNN model

Section 02 (R codes)

```
knitr::opts_chunk$set(echo = TRUE}  
  
----eval=FALSE, include=FALSE-----  
## Question 01  
  
# 1a)  
  
max_features <- 5000 # or 5000  
imdb <- dataset_imdb(num_words = max_features)  
c(c(x_train, y_train), c(x_test, y_test)) %<-% imdb  
  
# a function to decode a review  
word_index <- dataset_imdb_word_index()  
decode_review <- function(text, word_index) {  
  word <- names(word_index)  
  idx <- unlist(word_index, use.names = FALSE)  
  word <- c("<PAD>", "<START>", "<UNK>", "<UNUSED>", word)  
  idx <- c(0:3, idx + 3)  
  words <- word[match(text, idx, 2)]  
  paste(words, collapse = " ")  
}  
  
library(Matrix)  
one_hot <- function(sequences, dimension) {  
  seqlen <- sapply(sequences, length)  
  n <- length(seqlen)  
  rowind <- rep(1:n, seqlen)  
  colind <- unlist(sequences)  
  sparseMatrix(i = rowind, j = colind,  
    dims = c(n, dimension))  
}  
  
x_train_1h <- one_hot(x_train, max_features)  
x_test_1h <- one_hot(x_test, max_features)  
  
# create a validation set of size 2000 leaving 23000 for training  
set.seed(3)  
ival <- sample(seq(along = y_train), 2000)  
  
### restrict the document lengths to the last L = 500 words  
maxlen <- 500  
x_train <- pad_sequences(x_train, maxlen = maxlen)  
x_test <- pad_sequences(x_test, maxlen = maxlen)  
  
----eval=FALSE, include=FALSE-----  
model <- keras_model_sequential() %>%  
  layer_embedding(input_dim = max_features, output_dim = 32) %>%  
  layer_lstm(units = 32) %>%  
  layer_dense(units = 1, activation = "sigmoid")  
  
model %>% compile(optimizer = "rmsprop",  
  loss = "binary_crossentropy", metrics = c("acc"))  
history <- model %>% fit(x_train, y_train, epochs = 15,  
  batch_size = 128, validation_data = list(x_test, y_test))  
  
plot(history)  
predy <- predict(model, x_test) > 0.5  
mean(abs(y_test == as.numeric(predy)))
```

```

-----eval=FALSE, include=FALSE-----
# 1b)

max_features <- 5000 # or 5000
imdb <- dataset_imdb(num_words = max_features)
c(c(x_train, y_train), c(x_test, y_test)) %<-% imdb

# a function to decode a review
word_index <- dataset_imdb_word_index()
decode_review <- function(text, word_index) {
  word <- names(word_index)
  idx <- unlist(word_index, use.names = FALSE)
  word <- c("<PAD>", "<START>", "<UNK>", "<UNUSED>", word)
  idx <- c(0:3, idx + 3)
  words <- word[match(text, idx, 2)]
  paste(words, collapse = " ")
}

library(Matrix)
one_hot <- function(sequences, dimension) {
  seqlen <- sapply(sequences, length)
  n <- length(seqlen)
  rowind <- rep(1:n, seqlen)
  colind <- unlist(sequences)
  sparseMatrix(i = rowind, j = colind,
    dims = c(n, dimension))
}

x_train_1h <- one_hot(x_train, max_features)
x_test_1h <- one_hot(x_test, max_features)

# create a validation set of size 2000 leaving 23000 for training
set.seed(3)
ival <- sample(seq(along = y_train), 2000)

### restrict the document lengths to the last L = 500 words
maxlen <- 500
x_train <- pad_sequences(x_train, maxlen = maxlen)
x_test <- pad_sequences(x_test, maxlen = maxlen)

-----eval=FALSE, include=FALSE-----
model <- keras_model_sequential() %>%
  layer_embedding(input_dim = max_features, output_dim = 32) %>%
  layer_lstm(units = 32) %>%
  layer_dense(units = 64, activation = "relu") %>%
  layer_dense(units = 32, activation = "relu") %>%
  layer_dense(units = 1, activation = "sigmoid")

model %>% compile(optimizer = "rmsprop",
  loss = "binary_crossentropy", metrics = c("acc"))
history <- model %>% fit(x_train, y_train, epochs = 15,
  batch_size = 128, validation_data = list(x_test, y_test))

plot(history)
predy <- predict(model, x_test) > 0.5
mean(abs(y_test == as.numeric(predy)))

```

```

----eval=FALSE, include=FALSE-----
#####
## Question 02

# Data Pre processing

#set up the data, and standardize each of the variables
library(ISLR2)
xdata <- data.matrix(
  NYSE[, c("DJ_return", "log_volume","log_volatility")]
)
istrain <- NYSE[, "train"]
xdata <- scale(xdata)

#functions to create lagged versions of the three time series
lagm <- function(x, k = 1) {
  n <- nrow(x)
  pad <- matrix(NA, k, ncol(x))
  rbind(pad, x[1:(n - k), ])
}

#use this function to create a data frame with all the required lags, as well as the response variable.
arframe <- data.frame(log_volume = xdata[, "log_volume"],
  L1 = lagm(xdata, 1), L2 = lagm(xdata, 2),
  L3 = lagm(xdata, 3), L4 = lagm(xdata, 4),
  L5 = lagm(xdata, 5)
)

#remove NA rows, and adjust istrain accordingly
arframe <- arframe[-(1:5), ]
istrain <- istrain[-(1:5)]

----eval=FALSE, include=FALSE-----
# 2a) fit an RNN
# reshape data to fit RNN
n <- nrow(arframe)
xrnn <- data.matrix(arframe[, -1])
xrnn <- array(xrnn, c(n, 3, 5))
xrnn <- xrnn[, , 5:1]
xrnn <- aperm(xrnn, c(1, 3, 2))

model <- keras_model_sequential() %>%
  layer_simple_rnn(units = 12,
    input_shape = list(5, 3),
    dropout = 0.1, recurrent_dropout = 0.1) %>%
  layer_dense(units = 1)
model %>% compile(optimizer = optimizer_rmsprop(),
  loss = "mse")

history <- model %>% fit(
  xrnn[istrain,, ], arframe[istrain, "log_volume"],
  batch_size = 64, epochs = 100,
  validation_data =
    list(xrnn[!istrain,, ], arframe[!istrain, "log_volume"])
)

kpred <- predict(model, xrnn[!istrain,, ])
V0 <- var(arframe[!istrain, "log_volume"])
1 - mean((kpred - arframe[!istrain, "log_volume"])^2) / V0

----eval=FALSE, include=FALSE-----

```

```

# 2a) fit an RNN with day_of_week added

# to include day_of_week to arframe
arframed <- data.frame(day = NYSE [-(1:5), "day_of_week"], arframe)

# reshape data to fit RNN
n <- nrow(arframed)
xrnn <- data.matrix(arframed[, -1])
xrnn <- array(xrnn, c(n, 3, 5))
xrnn <- xrnn[, , 5:1]
xrnn <- aperm(xrnn, c(1, 3, 2))

model <- keras_model_sequential() %>%
  layer_simple_rnn(units = 12,
    input_shape = list(5, 3),
    dropout = 0.1, recurrent_dropout = 0.1) %>%
  layer_dense(units = 1)
model %>% compile(optimizer = optimizer_rmsprop(),
  loss = "mse")

history <- model %>% fit(
  xrnn[istrain, , ], arframed[istrain, "log_volume"],
  batch_size = 64, epochs = 100,
  validation_data =
    list(xrnn[!istrain, , ], arframed[!istrain, "log_volume"])
)

kpred <- predict(model, xrnn[!istrain, , ])
V0 <- var(arframed[!istrain, "log_volume"])
1 - mean((kpred - arframed[!istrain, "log_volume"])^2) / V0

-----eval=FALSE, include=FALSE-----
# 2b) fit an RNN

# reshape data to fit RNN
n <- nrow(arframe)
xrnn <- data.matrix(arframe[, -1])
xrnn <- array(xrnn, c(n, 3, 5))
xrnn <- xrnn[, , 5:1]
xrnn <- aperm(xrnn, c(1, 3, 2))

model <- keras_model_sequential() %>%
  layer_simple_rnn(units = 12,
    input_shape = list(5, 3),
    dropout = 0.1, recurrent_dropout = 0.1) %>%
  layer_dense(units = 32, activation = "relu") %>%
  layer_dense(units = 1)
model %>% compile(optimizer = optimizer_rmsprop(),
  loss = "mse")

history <- model %>% fit(
  xrnn[istrain, , ], arframe[istrain, "log_volume"],
  batch_size = 64, epochs = 100,
  validation_data =
    list(xrnn[!istrain, , ], arframe[!istrain, "log_volume"])
)

kpred <- predict(model, xrnn[!istrain, , ])
V0 <- var(arframe[!istrain, "log_volume"])
1 - mean((kpred - arframe[!istrain, "log_volume"])^2) / V0

```