

# Deep Learning

Indrajith Wasala Mudiyanse

## Section 1

### Question 01

- a) Since the dataset is large, fitting multinomial logistic regression (MLR) using the “multinom” function is difficult. Therefore I used shallow learning to fit MLR. For this model, the training error is 0.2502114 and the validation error is 0.2799808.
- b) The model with two hidden layers is the best of the models considered.

Model	Training error	Validation error
NNM with 1 hidden layer	0.05511252	0.1085238
NNM with 2 hidden layers	0.0254432	0.09725994
NNM with 3 hidden layers	0.03340233	0.1080917

Table 1: Tabular summary of errors

- c) Among the used grid of parameters the best regularization parameter is 0.0001. The corresponding validation error for this parameter is 0.1265614. Therefore, the regularization does not help ( $0.1265614 > 0.09725994$ ).
- d) Among the used parameters the best dropout rate is 0.25. The corresponding validation error for this parameter is 0.1117960. Therefore, the dropout rate does not help.
- e) Since the validation errors in parts c and d do not improve, I recommend the best model in part b (the basic model with two hidden layers). The test error is 0.1406495 (training error= 0.0254432 and validation error= 0.09725994).
- f) To achieve the best results, increasing the power of a neural network requires a combination of techniques, experimentation, and tuning. For example, increase the network size, use a larger training dataset, use better optimization algorithms, and so on.

### Question 02

- a) The training error is 3.308168 and the validation error is 3.517034.
- b) The model with one hidden layers is the best of the models considered.

Model	Training error	Validation error
NNM with 1 hidden layer	1.720874	2.360861
NNM with 2 hidden layers	1.551988	2.415809

Table 2: Tabular summary of errors

- c) Among the used grid of parameters the best regularization parameter is 0.0005. The corresponding validation error for this parameter is 2.118346. Therefore, the regularization improves the model ( $2.118346 < 2.360861$ ).
- d) Among the used parameters the best dropout rate is 0.25. The corresponding validation error for this parameter is 2.221878. Therefore, the dropout rate improves the model (But not like regularization).
- e) I recommend the model in part c (the model with one hidden layers and regularization parameter 0.0005). The test error is 2.815951 (training error= 1.924338 and validation error= 2.118346).
- f) To achieve the best results, increasing the power of a neural network requires a combination of techniques, experimentation, and tuning. For example, increase the network size, use a larger training dataset, use better optimization algorithms, and so on.

## Section 2 (R codes)

```
knitr::opts_chunk$set(echo = TRUE}  
  
## ----include=FALSE-----  
## Question 01  
  
library(keras)  
  
mnist <- dataset_mnist()  
train_images <- mnist$train$x  
train_labels <- mnist$train$y  
test_images <- mnist$test$x  
test_labels <- mnist$test$y  
  
### Preprocess so that the training and test data for features are  
### in matrix form (as needed for model fitting)  
  
train_images <- array_reshape(train_images, c(60000, 28*28)) # matrix  
train_images <- train_images/255 # ensures all values are in [0, 1]  
test_images <- array_reshape(test_images, c(10000, 28*28))  
test_images <- test_images/255  
  
### Make the response categorical  
  
train_labels <- to_categorical(train_labels)  
test_labels <- to_categorical(test_labels)  
  
# Validation data split (80-20)  
set.seed(6390)  
rand<-sample(1:60000,size = 12000) # 60000*0.2=12000  
  
validation_images<-train_images[rand,]  
validation_labels<-train_labels[rand,]  
  
new_train_images<-train_images[-rand,]  
new_train_labels<-train_labels[-rand,]  
  
## ----include=FALSE-----  
## # a)  
### set up the network architecture  
  
network <- keras_model_sequential() %>%  
  layer_dense(units = 10, activation = "softmax", input_shape = c(28*28))  
  
### Compile the model  
  
network %>% compile(  
  optimizer = "rmsprop",  
  loss = "categorical_crossentropy", # loss function to minimize  
  metrics = c("accuracy") # monitor classification accuracy  
)  
  
network %>% fit(new_train_images, new_train_labels, epochs = 15,  
  batch_size = 128, validation_data = list(validation_images, validation_labels))  
  
### Examine training and validation performance  
metrics_train <- network %>% evaluate(new_train_images, new_train_labels)
```

```

metrics_val <- network %>% evaluate(validation_images,validation_labels)

## ----include=FALSE-----
## # b)
##
### set up the network architectures

# Model with one hidden layer
network_b1 <-keras_model_sequential() %>%
  layer_dense(units = 64, activation = "relu", input_shape = c(28*28)) %>%
  layer_dense(units = 10, activation = "softmax")

# Model with two hidden layers
network_b2 <-keras_model_sequential() %>%
  layer_dense(units = 64, activation = "relu", input_shape = c(28*28)) %>%
  layer_dense(units = 128, activation = "relu") %>%
  layer_dense(units = 10, activation = "softmax")

# Model with three hidden layers
network_b3 <-keras_model_sequential() %>%
  layer_dense(units = 64, activation = "relu", input_shape = c(28*28)) %>%
  layer_dense(units = 128, activation = "relu") %>%
  layer_dense(units = 256, activation = "relu") %>%
  layer_dense(units = 10, activation = "softmax")

### Compile the models
network_b1 %>% compile(
  optimizer = "rmsprop",
  loss = "categorical_crossentropy", # loss function to minimize
  metrics = c("accuracy") # monitor classification accuracy
)

network_b2 %>% compile(
  optimizer = "rmsprop",
  loss = "categorical_crossentropy", # loss function to minimize
  metrics = c("accuracy") # monitor classification accuracy
)

network_b3 %>% compile(
  optimizer = "rmsprop",
  loss = "categorical_crossentropy", # loss function to minimize
  metrics = c("accuracy") # monitor classification accuracy
)

### Fit the models

network_b1 %>% fit(new_train_images, new_train_labels, epochs = 10,
  batch_size = 128,validation_data = list(validation_images,validation_labels))
network_b2 %>% fit(new_train_images, new_train_labels, epochs = 10,
  batch_size = 128,validation_data = list(validation_images,validation_labels))
network_b3 %>% fit(new_train_images, new_train_labels, epochs = 10,
  batch_size = 128,validation_data = list(validation_images,validation_labels))

### Examine training and validation performance
metrics_train_b1 <- network_b1 %>% evaluate(new_train_images, new_train_labels)
metrics_val_b1 <- network_b1 %>% evaluate(validation_images, validation_labels)

metrics_train_b2 <- network_b2 %>% evaluate(new_train_images, new_train_labels)
metrics_val_b2 <- network_b2 %>% evaluate(validation_images, validation_labels)

```

```

metrics_train_b3 <- network_b3 %>% evaluate(train_images, train_labels)
metrics_val_b3 <- network_b3 %>% evaluate(validation_images, validation_labels)

## ----include=FALSE-----
## # c)
##
# Add L2 weight regularization to the best model

para_grid<-c(0.0001,0.0005,0.001)
loss_c<-c()

for (i in 1:length(para_grid)) {

  best_model_reg <- keras_model_sequential() %>%
    layer_dense(units = 64, activation = "relu", input_shape = c(28*28),
      kernel_regularizer = regularizer_l2(para_grid[i])) %>%
    layer_dense(units = 128, activation = "relu",
      kernel_regularizer = regularizer_l2(para_grid[i])) %>%
    layer_dense(units = 10, activation = "softmax")

  best_model_reg %>% compile(
    optimizer = "rmsprop",
    loss = "categorical_crossentropy", # loss function to minimize
    metrics = c("accuracy") # monitor classification accuracy
  )

  best_model_reg %>% fit(new_train_images, new_train_labels, epochs = 10,
    batch_size = 128,
    validation_data=list(validation_images,validation_labels))
  metrics_val_best <- best_model_reg %>% evaluate(validation_images,validation_labels)
  loss_c[i]<- metrics_val_best[1]
}

## ----include=FALSE-----
## # d)
drop_grid<-c(0.25,0.50)
loss_d<-c()

for (i in 1:length(drop_grid)) {

  best_model_reg <- keras_model_sequential() %>%
    layer_dense(units = 64, activation = "relu", input_shape = c(28*28)) %>%
    layer_dropout(rate = drop_grid[i]) %>%
    layer_dense(units = 128, activation = "relu") %>%
    layer_dropout(rate = drop_grid[i]) %>%
    layer_dense(units = 10, activation = "softmax")

  best_model_reg %>% compile(
    optimizer = "rmsprop",
    loss = "categorical_crossentropy", # loss function to minimize
    metrics = c("accuracy") # monitor classification accuracy
  )

  best_model_reg %>% fit(new_train_images, new_train_labels, epochs = 15,
    batch_size = 128,validation_data = list(validation_images,validation_labels))

  metrics_val_best <- best_model_reg %>% evaluate(validation_images,validation_labels)

```

```

    loss_d[i]<- metrics_val_best[1]
}

## ----include=FALSE-----
## # e)
## # Calculate test error

network_b2 %>% fit(train_images, train_labels, epochs = 10, batch_size = 128)

metrics_test_b2 <- network_b2 %>% evaluate(test_images, test_labels)

#####
## ----include=FALSE-----
## Question 02

library(keras)
boston <- dataset_boston_housing()
c(c(train_data, train_targets), c(test_data, test_targets)) %<-% boston

### Preprocess the data

### Standardize the training and test features

mean <- apply(train_data, 2, mean)
std <- apply(train_data, 2, sd)
train_data <- scale(train_data, center = mean, scale = std)
test_data <- scale(test_data, center = mean, scale = std)

## ----include=FALSE-----
## # a)

library(glmnet)
new_train <- unlist(train_data)
new_train <- matrix(new_train, nrow = 404, ncol = 13)
new_target <- as.vector(unlist(train_targets))
cvfit <- cv.glmnet(new_train, new_target,
  type.measure = "mae", nfolds = 4)

trn_pred <- predict(cvfit, newx = new_train, s = "lambda.min")

# Calculate training and cross-validation errors
train_error <- mean(abs(new_target - trn_pred))
cv_error <- cvfit$cvm[which.min(cvfit$cvm)]

## ----include=FALSE-----
## # b)

### set up the network architectures

# Model with one hidden layer
Q2network_b1 <-keras_model_sequential() %>%
  layer_dense(units = 64, activation = "relu", input_shape = c(13)) %>%
  layer_dense(units = 1, activation = "linear")

# Model with two hidden layers
Q2network_b2 <-keras_model_sequential() %>%
  layer_dense(units = 64, activation = "relu", input_shape = c(13)) %>%

```

```

layer_dense(units = 128, activation = "relu") %>%
layer_dense(units = 1, activation = "linear")

### Compile the models
Q2network_b1 %>% compile(
  optimizer = "rmsprop",
  loss = "mse",      # loss function to minimize
  metrics = c("mae")
)

Q2network_b2 %>% compile(
  optimizer = "rmsprop",
  loss = "mse",      # loss function to minimize
  metrics = c("mae")
)

k<-4
set.seed(1)
indices <- sample(1:nrow(train_data),replace = FALSE)
folds <- cut(indices, breaks = k, labels = FALSE)

Q2b1_train<-c()
Q2b1_val<-c()

Q2b2_train<-c()
Q2b2_val<-c()

for (i in 1:k) {

  val_ind<- which(folds == i, arr.ind = TRUE)
  val_x<-train_data[val_ind,]
  train_x<-train_data[-val_ind,]

  train_y<-train_targets[-val_ind]
  val_y<-train_targets[val_ind]

  Q2network_b1 %>% fit(train_x, train_y, epochs = 100,
                        batch_size = 16)
  Q2network_b2 %>% fit(train_x, train_y, epochs = 100,
                        batch_size = 16)

  ### Examine training and validation performance
  metrics_train_Q2b1 <- Q2network_b1 %>% evaluate(train_x, train_y)
  metrics_val_Q2b1 <- Q2network_b1 %>% evaluate(val_x, val_y)
  Q2b1_train[i]<-metrics_train_Q2b1[2]
  Q2b1_val[i]<-metrics_val_Q2b1[2]

  metrics_train_Q2b2 <- Q2network_b2 %>% evaluate(train_x, train_y)
  metrics_val_Q2b2 <- Q2network_b2 %>% evaluate(val_x, val_y)
  Q2b2_train[i]<-metrics_train_Q2b2[2]
  Q2b2_val[i]<-metrics_val_Q2b2[2]

}

# Model 1 training error and validation error
error_Q2b1_train<- mean(Q2b1_train)
error_Q2b1_val<-mean(Q2b1_val)

# Model 2 training error and validation error
error_Q2b2_train<- mean(Q2b2_train)
error_Q2b2_val<-mean(Q2b2_val)

```

```

## ----include=FALSE-----
## # c)

# Add L2 weight regularization to the best model

para_grid<-c(0.0001,0.0005,0.001)
k<-4
  set.seed(1)
  indices <- sample(1:nrow(train_data),replace = FALSE)
  folds <- cut(indices, breaks = k, labels = FALSE)

  Q3_train<-c()
  Q3_val<-c()

  l1_erorr<-c()

for (j in 1:length(para_grid)) {
  # Model with one hidden layer
Q3network <-keras_model_sequential() %>%
  layer_dense(units = 64, activation = "relu",
              input_shape = c(13),kernel_regularizer = regularizer_l2(para_grid[j])) %>%
  layer_dense(units = 1, activation = "linear")

Q3network %>% compile(
  optimizer = "rmsprop",
  loss = "mse",      # loss function to minimize
  metrics = c("mae")
)

  for (i in 1:k) {

    val_ind<- which(folds == i, arr.ind = TRUE)
    val_x<-train_data[val_ind,]
    train_x<-train_data[-val_ind,]

    train_y<-train_targets[-val_ind]
    val_y<-train_targets[val_ind]

    Q3network %>% fit(train_x, train_y, epochs = 100,
                     batch_size = 16)

    ### Examine training and validation performance
    metrics_train_Q3 <- Q3network %>% evaluate(train_x, train_y)
    metrics_val_Q3 <- Q3network %>% evaluate(val_x, val_y)
    Q3_train[i]<-metrics_train_Q3[2]
    Q3_val[i]<-metrics_val_Q3[2]

  }
  l1_erorr[j]<-mean(Q3_val)
}

## ----include=FALSE-----
## # d)

drop_grid<-c(0.25,0.50)
k<-4

```

```

set.seed(1)
indices <- sample(1:nrow(train_data),replace = FALSE)
folds <- cut(indices, breaks = k, labels = FALSE)

Q2d_train<-c()
Q2d_val<-c()

D0_erorr<-c()

for (j in 1:length(drop_grid)) {
  # Model with one hidden layer
  Q2d_network <-keras_model_sequential() %>%
  layer_dense(units = 64, activation = "relu", input_shape = c(13)) %>%
  layer_dropout(rate = drop_grid[j]) %>%
  layer_dense(units = 1, activation = "linear")

  Q2d_network %>% compile(
    optimizer = "rmsprop",
    loss = "mse",      # loss function to minimize
    metrics = c("mae")
  )

  for (i in 1:k) {

    val_ind<- which(folds == i, arr.ind = TRUE)
    val_x<-train_data[val_ind,]
    train_x<-train_data[-val_ind,]

    train_y<-train_targets[-val_ind]
    val_y<-train_targets[val_ind]

    Q2d_network %>% fit(train_x, train_y, epochs = 100,
                       batch_size = 16)

    ### Examine training and validation performance
    metrics_train_Q2d <- Q2d_network %>% evaluate(train_x, train_y)
    metrics_val_Q2d <- Q2d_network %>% evaluate(val_x, val_y)
    Q2d_train[i]<-metrics_train_Q2d[2]
    Q2d_val[i]<-metrics_val_Q2d[2]

  }
  D0_erorr[j]<-mean(Q2d_val)
}

## ----include=FALSE-----
## # e)

# Model with one hidden layer
Q2e_network <-keras_model_sequential() %>%
  layer_dense(units = 64, activation = "relu",
              input_shape = c(13),kernel_regularizer = regularizer_l2(0.0005)) %>%
  layer_dense(units = 1, activation = "linear")

Q2e_network %>% compile(
  optimizer = "rmsprop",
  loss = "mse",      # loss function to minimize
  metrics = c("mae")

```



)

```
Q2e_network %>% fit(train_data, train_targets, epochs = 100,  
                    batch_size = 16)
```

```
### Examine training and validation performance
```

```
metrics_train_Q2e <- Q2e_network %>% evaluate(train_data, train_targets)
```

```
metrics_val_Q2e <- Q2e_network %>% evaluate(test_data, test_targets)
```